

Comprehensive End-to-End Run — v6 Vocabulary, Real Loan Data, Full Audit Trail

Run: `run_013_comprehensive_e2e_v6` · **Date:** 2026-07-27 · **Cost:** \$0.00 (100% deterministic — no LLM calls anywhere in this run) · **Test suite:** 318/318 passing

This is the genuine end-to-end proof you asked for after `run_012` (compile-time coverage only): the 5 real loans, run through the actual engine, with real derived loan-side data, replayed against the deployed baseline, with every stage logged for audit.

Part 1 — Scope: what's "real" in this run

`run_012` proved 1,530 of 3,203 checks now *compile* with a real precondition. This run proves what actually *happens* when those checks execute against real loans — which requires the loan side of the equation too, not just the AMQ side.

Of the vocabulary's 16 facts, only **3 have a genuine, already-extracted signal** on the 5 real synthetic loans (checked field-by-field before writing any derivation code — see `build_loan_profiles_v2.py` 's docstring):

Fact	Derived from	Coverage
<code>gift_funds_used</code>	<code>doc_present_gift_letter</code>	5/5 loans (v1, unchanged)
<code>loan_transaction_type</code>	<code>loan_purpose_general_1003</code>	5/5 loans
<code>appraisal_in_file</code>	<code>doc_present_va_appraisal /</code> <code>doc_present_usda_appraisal /</code> <code>appraised_value</code>	4/5 loans — <code>loan_02</code> (FHA) is honestly underivable, not defaulted to false

The other 13 facts have no direct signal in these fixtures (most are deliberately narrow — `loan_05` has 28 extracted fields total vs. `loan_01`'s 216 — so a missing field means "not extracted for this fixture," not "doesn't exist"). Deriving them would mean inventing an inference from adjacent data, the same category of mistake the citation-mismatch fix just corrected. **This run is honest about that scope — it proves the pipeline works end-to-end for the facts that are genuinely wired, not the full 16.**

Part 2 — The 5-loan run

Loan	Disposition	FAIL	NEEDS_REVIEW	NOT_APPLICABLE	PASS
2025-0917-001 (Conventional Purchase)	NEEDS_REVIEW	1,643	1,011	536	13
2025-1004-FHA-002 (FHA Purchase)	NEEDS_REVIEW	1,384	1,310	506	3
2025-1108-VA-003 (VA Purchase)	NEEDS_REVIEW	1,646	1,010	544	3

2025-1215-FRD-004 (Freddie Mac Cash-Out Refi)	NEEDS_REVIEW	1,639	1,008	547	9
2025-1122-USDA-005 (USDA RHS 502)	NEEDS_REVIEW	1,650	1,006	544	3

Every loan still lands on NEEDS_REVIEW at this scale (this is the 3,203-check Retail rulebook, not the smaller known-defects set from earlier reports) — expected, since most checks still can't resolve applicability without the 13 not-yet-derivable facts.

Part 3 — Replay vs. the deployed baseline: 5,956 flips

Transition	Count	Read
FAIL → NEEDS_REVIEW	4,278	The dominant transition. The deployed baseline had no precondition wiring — it asserted FAIL on thousands of checks whose applicability it couldn't actually confirm. The new system honestly abstains instead.
FAIL → NOT_APPLICABLE	1,142	Clean win — precondition confirmed false, correctly gated off.
NOT_APPLICABLE → NEEDS_REVIEW	448	The one direction worth a caveat — see below.
NEEDS_REVIEW → NOT_APPLICABLE	82	Clean win — an unknown resolved to a confident non-match.
PASS → NEEDS_REVIEW	2	Spot-checked individually — see Part 4.
PASS → NOT_APPLICABLE	4	Spot-checked individually — see Part 4.

On the 448 NOT_APPLICABLE → NEEDS_REVIEW cases: the deployed baseline reached NOT_APPLICABLE through a different, kind-specific path (e.g. "field is absent → not applicable"), not through precondition logic it never had. Now that a real (but unresolvable) precondition is attached to these checks, they correctly surface as "we don't know" rather than reusing that older, narrower path. This is a legitimate, disclosed trade — not silently swept under "good news only."

Part 4 — The 6 cases that moved *away* from PASS (verified individually, not assumed safe)

These are the one direction that would look like a real regression, so each was checked by hand against what the check and loan actually are:

Check	Loan	Old → New	Verdict
-------	------	-----------	---------

fha-203h-min-credit-score-500	Conventional Purchase	PASS → NEEDS_REVIEW	✅ Improvement — this is an FHA disaster-relief-program check; the loan is Conventional. The old PASS was incidental (nothing gated it from running at all); the new NEEDS_REVIEW honestly says "can't confirm this program applies," which is correct — it's not confirmed NOT to apply either without more data.
homereday-sweat-equity-ltv-max	Conventional Purchase, Freddie Mac Cash-Out Refi	PASS → NOT_APPLICABLE	✅ Improvement — a HomeReady-specific check; neither loan is HomeReady. NOT_APPLICABLE is the <i>more correct</i> verdict than a PASS that was never really evaluating anything relevant.
o-fnm-00531-fnm-commitment-present	Conventional Purchase	PASS → NEEDS_REVIEW	✅ Improvement — same pattern: an investor-commitment check that previously ran unconditionally.
va-type1-refi-fixed-to-arm-ltv-discount-point	Conventional Purchase, Freddie Mac Cash-Out Refi	PASS → NOT_APPLICABLE	✅ Improvement — a VA fixed-to-ARM refinance check. Neither loan is VA, and one of them (Conventional Purchase) isn't even a refinance. NOT_APPLICABLE is unambiguously correct here.

All 6 are correctness improvements, not regressions — every one replaces a check that was running (and passing) with no real basis for applying, with the correct "doesn't apply" or "can't confirm" verdict.

Part 5 — Full audit trail

Every stage of this run is logged to `storage/logs/run_013_comprehensive_e2e_v6.jsonl` — **24,050 structured events**, per this project's own evidence-chain requirement (input → method → verdict, no black boxes):

Stage	Events	What's logged
loan_profile_derivation	15	Every derived/underivable fact, per loan, with its exact source field/value
precondition_attachment	2,052	Every one of the 1,530 attached + 520 flagged checks, with the resolved condition or refusal reason
engine_execution	16,020	Every check, every loan — input, evaluation method, verdict
replay	5,957	Every status flip vs. the deployed baseline
cost	1	Explicit \$0.00 / 0 LLM calls / 100% deterministic-resolution — reported, not just omitted

Nothing in this run is a summary-only claim — every number above traces to a specific, inspectable log line.

Part 6 — What this proves, and what it doesn't

Proves: the compiled ruleset + expanded vocabulary + engine, wired together, execute correctly end-to-end against real loan documents, for the facts genuinely available today. The precondition mechanism itself is validated as correct (all 6 PASS-departures checked by hand are improvements). Zero LLM cost, fully deterministic, fully logged.

Doesn't prove: that the other 13 vocabulary facts work correctly against real loans — they have no loan-side data yet, so every check gated on them honestly surfaces as NEEDS_REVIEW rather than a real disposition. Closing that gap is real extraction work (Non-Negotiable #2 — Touchless's contract), tracked separately, not something this run can shortcut.

Where everything lives

Artifact	Path
This run's script	p0/compile_runs/run_013_comprehensive_e2e_v6/build_and_run.py
Full results (JSON)	result/qc_results/run_013_comprehensive_e2e_v6_results.json
Compiled ruleset (unsigned)	result/rules/comprehensive_e2e_v6_ruleset.json
Full audit log (24,050 events)	storage/logs/run_013_comprehensive_e2e_v6.jsonl
Loan profiles (v2, 3 facts derived)	storage/loan_profiles/v2/loan_0N.json
Derivation logic + feasibility writeup	p0/qc_engine/build_loan_profiles_v2.py
Structured logging module	p0/qc_engine/eval_log.py
Tests	test_loan_profiles_v2.py, test_eval_log.py (20 new tests, 318 total passing)